

Scheduling Notes

Metrics Calculation

Throughput

N = jobs completed

T = time to complete all jobs

$$\frac{N}{T}$$

Avg. Turnaround Time

C = individual job completion time

N = jobs completed

$$\frac{\sum C}{N}$$

Avg. Wait Time

C = individual job wait time

N = jobs completed

$$\frac{\sum C}{N}$$

Avg. Response Time


C = individual job response time == first time process begins execution

N = jobs completed

$$\frac{\sum C}{N}$$

FCFS


Task	Time Units
A	8
B	4
C	9
D	5



Metric	FCFS
Turn around time	$(8+12+21+26)/4 = 16.5$ sec/process
Waiting	$(0+8+12+21)/4 = 10.25$ sec/process
Throughput	$4/(26) = .15$ process/sec
Response	$(0+8+12+21)/4 = 10.25$ sec/process

SJF

Task	Time Units
A	8
B	4
C	9
D	5




--	--

Metric	FCFS
Turn around time	$(4+9+17+26)/4 = 14 \text{ sec/process}$
Waiting	$(0+4+9+17)/4 = 7.5 \text{ sec/process}$
Throughput	$4/(26) = .15 \text{ process/sec}$
Response	$(0+4+9+17)/4 = 7.5 \text{ sec/process}$

RR

Task	Time Units
A	8
B	4
C	9
D	5

Time Quanta (time slice) == 3



Metric	FCFS
Turn around time	$(23+16+26+21)/4 = 21.5 \text{ sec/process}$
Waiting	$((21-3-3)+(15-3-3)+(23-3-3)+(19-3-3))/4 = 15 \text{ sec/process}$
Throughput	$4/(26) = .15 \text{ process/sec}$
Response	$(0+3+6+9)/4 = 4.5 \text{ sec/process}$

Note for round robin wait time we have to subtract the execution time. The easiest way to do this is find the final time slice and subtract t at that time from the number of time slices execute * the time quanta. So you could rewrite the equation

$$\frac{\sum t_{at_last_slice} - \#_of_execution_slices * time_slice}{\#_of_processes}$$

$$\frac{\sum (21 - 2 * 3) + (15 - 1 * 3) + (23 - 2 * 3) + (19 - 1 * 3)}{4}$$

Timeslice

CPU Bound Tasks: *prefer longer timeslices*

- Limits context switching overheads
- Keeps CPU utilization and throughput high

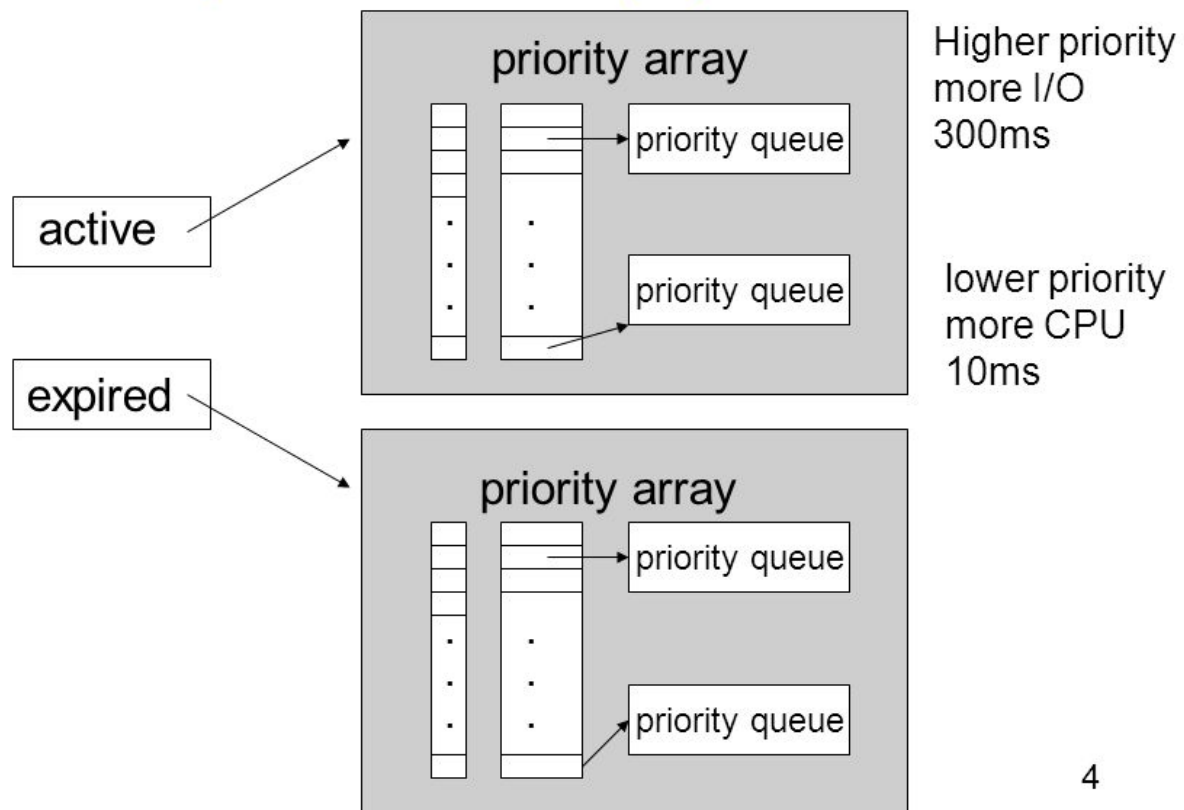
I/O Bound Tasks: *prefer shorter timeslices*

- I/O bound tasks can issue I/O ops earlier
 - Keeps CPU and device utilization high
 - Better user-perceived performance
-

Linux O(1) Scheduler

1. $O(1)$ == constant time to select/add task to the run queue regardless of task count
2. Preemptive, priority-based
 - realtime (0-99)
 - timesharing (100-139)
3. User Processes
 - default == 120
 - nice value(-20 to 19)
 - system call to span 100-139
4. Borrows two traits from MLFQ, though these are implemented differently than a pure MLFQ
 - *Timeslice*: the time quantum is associated with the priority level
 - depends on priority
 - smallest for low priority
 - highest for high priority
 - *Feedback*: maintains knowledge of the "history" of the tasks
 - sleep time: waiting/idling
 - longer sleep means interactive task- priority **-5** (boost)
 - smaller sleep mean CPU intensive - priority **+5** (lowered)
 - these will increment and decrement between 100-139

Runqueue for O(1) Scheduler



4

5. Runqueue (ready queue) - There are **2 arrays** of tasks

- **Active:**
 - used to pick the next task to run
 - constant time to add/select
 - uses a bitmap to access array
 - constant time to add and pull from queue at array index
- **Expired:**
 - inactive list: meaning task will not be scheduled as long as there is something to run in the active array
 - when no more tasks are in active array the pointers to the expired and active arrays swap
 - This helps ensure fairness, though it does not guarantee fairness
 - High to Low timeslice values associate with priority
 - The lower priority task are given low values which makes them more likely to expire
 - The higher priority task are given high timeslice values which makes them likely to stay in the active array

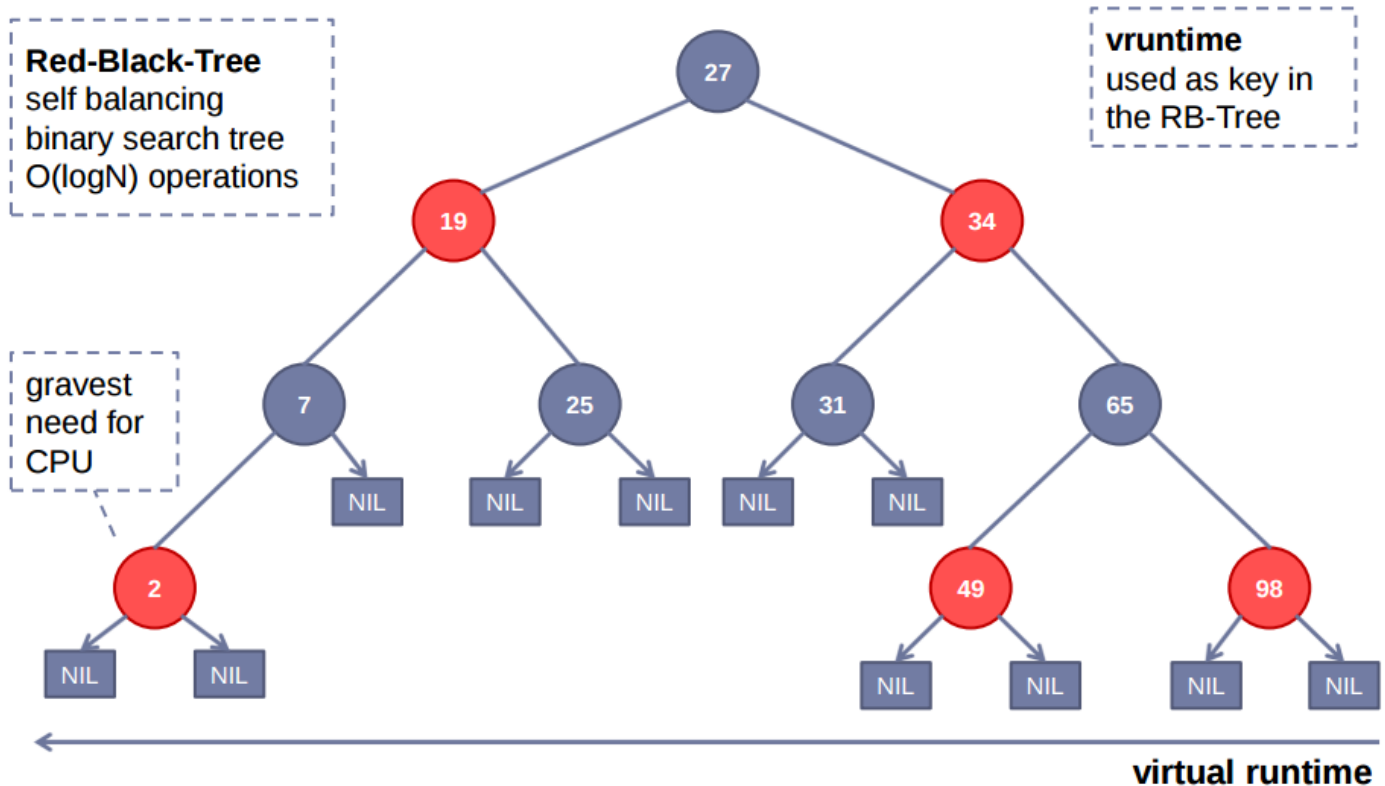
6. Performance

- poor performance of interactive task
 - when an interactive process is placed on the expired queue (timeslice expired), it will not be

given the CPU until **ALL** task on the active queue are removed

- This causes "jitter" in performance
- Does not guarantee fairness
 - starvation could occur (or poor performance) if there is always a process on the active queue

Linux CFS Scheduler (now default)



1. CFS was developed to address the issues with O(1) scheduler
2. Runqueue == Red-Black Tree
 - Self balancing
 - ordered by "vruntime"
 - vruntime == time spent on the cpu
 - tracked in a nanosecond granularity
3. Ordering
 - The left most child is the task that has spent the least amount of time on the CPU
 - The right most child is the task that has consumed more vruntime or CPU
4. CFS scheduling algorithm
 - always picks the left most node next
 - periodically adjust vruntime
 - will compare currently running task vruntime to the LMN in the tree
 - if smaller, continue running

- if larger, preempt and place in the tree
- runtime progress rate depends on priority of niceness
 - runtime progresses faster for low-priority
 - they will lose the cpu faster
 - runtime progress slower for high-priority task
 - they will keep the CPU longer
- uses one data structure (red black tree) for all priority levels

5. Performance

- select task => $O(1)$ time
- add task => $O(\log N)$
 - this is acceptable at current system loads, though it may be replaced as systems continue to support more task